



Deploy a Kubernetes Cluster with NVIDIA DeepOps Automated Deployment

NetApp Solutions

Dorian Henderson, Kevin Hoke
April 06, 2021

This PDF was generated from https://docs.netapp.com/us-en/netapp-solutions/ai/hc.ai_edge_deploy_a_kubernetes_cluster_with_nvidia_deepops_automated_deployment.htm on August 18, 2021. Always check docs.netapp.com for the latest.

Table of Contents

Deploy a Kubernetes Cluster with NVIDIA DeepOps Automated Deployment 1

Deploy a Kubernetes Cluster with NVIDIA DeepOps Automated Deployment

To deploy and configure the Kubernetes Cluster with NVIDIA DeepOps, complete the following steps:

1. Make sure that the same user account is present on all the Kubernetes master and worker nodes.
2. Clone the DeepOps repository.

```
git clone https://github.com/NVIDIA/deepops.git
```

3. Check out a recent release tag.

```
cd deepops  
git checkout tags/20.08
```

If this step is skipped, the latest development code is used, not an official release.

4. Prepare the Deployment Jump by installing the necessary prerequisites.

```
./scripts/setup.sh
```

5. Create and edit the Ansible inventory by opening a VI editor to `deepops/config/inventory`.
 - a. List all the master and worker nodes under [all].
 - b. List all the master nodes under [kube-master]
 - c. List all the master nodes under [etcd]
 - d. List all the worker nodes under [kube-node]

```
#####
# ALL NODES
# NOTE: Use existing hostnames here, DeepOps will conf
#####
[all]
hci-ai-k8-master-01      ansible_host=172.21.232.114
hci-ai-k8-master-02      ansible_host=172.21.232.115
hci-ai-k8-master-03      ansible_host=172.21.232.116
hci-ai-k8-worker-01      ansible_host=172.21.232.109
hci-ai-k8-worker-02      ansible_host=172.21.232.110

#####
# KUBERNETES
#####
[kube-master]
hci-ai-k8-master-01
hci-ai-k8-master-02
hci-ai-k8-master-03

# Odd number of nodes required
[etcd]
hci-ai-k8-master-01
hci-ai-k8-master-02
hci-ai-k8-master-03

# Also add mgmt/master nodes here if they will run non
[kube-node]
hci-ai-k8-worker-01
hci-ai-k8-worker-02

[k8s-cluster:children]
kube-master
kube-node
```

6. Enable GPUOperator by opening a VI editor to `deepops/config/group_vars/k8s-cluster.yml`.

```
# Provide option to use GPU Operator instead of setting up NVIDIA driver and
# Docker configuration.
deepops_gpu_operator_enabled: true
```

7. Set the value of `deepops_gpu_operator_enabled` to true.
8. Verify the permissions and network configuration.

```
ansible all -m raw -a "hostname" -k -K
```

- If SSH to the remote hosts requires a password, use -k.
- If sudo on the remote hosts requires a password, use -K.

9. If the previous step passed without any issues, proceed with the setup of Kubernetes.

```
ansible-playbook --limit k8s-cluster playbooks/k8s-cluster.yml -k -K
```

10. To verify the status of the Kubernetes nodes and the pods, run the following commands:

```
kubectl get nodes
```

```
rarvind@deployment-jump:~/deepops$ kubectl get nodes
NAME                                STATUS    ROLES    AGE    VERSION
hci-ai-k8-master-01                Ready    master   2d19h   v1.17.6
hci-ai-k8-master-02                Ready    master   2d19h   v1.17.6
hci-ai-k8-master-03                Ready    master   2d19h   v1.17.6
hci-ai-k8-worker-01                Ready    <none>    2d19h   v1.17.6
hci-ai-k8-worker-02                Ready    <none>    2d19h   v1.17.6
```

```
kubectl get pods -A
```

It can take a few minutes for all the pods to run.

```

rarvind@deployment-jump:~/deepops$ kubectl get pods -A
NAMESPACE      NAME                                                                 READY   STATUS
default        gpu-operator-74c97448d9-ppdlc                                       1/1     Running
default        nvidia-gpu-operator-node-feature-discovery-master-ffc5b57dx9wtl    1/1     Running
default        nvidia-gpu-operator-node-feature-discovery-worker-2lr9t            1/1     Running
default        nvidia-gpu-operator-node-feature-discovery-worker-6l6x7            1/1     Running
default        nvidia-gpu-operator-node-feature-discovery-worker-jf696            1/1     Running
default        nvidia-gpu-operator-node-feature-discovery-worker-tmtwv            1/1     Running
default        nvidia-gpu-operator-node-feature-discovery-worker-z4nlh            1/1     Running
gpu-operator-resources  nvidia-container-toolkit-daemonset-7jbl4                            1/1     Running
gpu-operator-resources  nvidia-container-toolkit-daemonset-x5ktb                            1/1     Running
gpu-operator-resources  nvidia-dcgm-exporter-5x94p                                          1/1     Running
gpu-operator-resources  nvidia-dcgm-exporter-7cbrl                                          1/1     Running
gpu-operator-resources  nvidia-device-plugin-daemonset-n8vrk                               1/1     Running
gpu-operator-resources  nvidia-device-plugin-daemonset-z7j6s                               1/1     Running
gpu-operator-resources  nvidia-device-plugin-validation                                     0/1     Completed
gpu-operator-resources  nvidia-driver-daemonset-7h752                                        1/1     Running
gpu-operator-resources  nvidia-driver-daemonset-v4rbj                                        1/1     Running
gpu-operator-resources  nvidia-driver-validation                                             0/1     Completed
kube-system        calico-kube-controllers-777478f4ff-jknxg                            1/1     Running
kube-system        calico-node-2j9mr                                                    1/1     Running
kube-system        calico-node-czk76                                                    1/1     Running
kube-system        calico-node-jpdxn                                                    1/1     Running
kube-system        calico-node-nwnvn                                                    1/1     Running
kube-system        calico-node-ssjrx                                                    1/1     Running
kube-system        coredns-76798d84dd-5pvqf                                            1/1     Running
kube-system        coredns-76798d84dd-w7l2j                                            1/1     Running
kube-system        dns-autoscaler-85f898cd5c-qqrbb                                     1/1     Running
kube-system        kube-apiserver-hci-ai-k8-master-01                                  1/1     Running
kube-system        kube-apiserver-hci-ai-k8-master-02                                  1/1     Running
kube-system        kube-apiserver-hci-ai-k8-master-03                                  1/1     Running
kube-system        kube-controller-manager-hci-ai-k8-master-01                        1/1     Running
kube-system        kube-controller-manager-hci-ai-k8-master-02                        1/1     Running
kube-system        kube-controller-manager-hci-ai-k8-master-03                        1/1     Running
kube-system        kube-proxy-5znxx                                                     1/1     Running
kube-system        kube-proxy-fk6h6                                                     1/1     Running
kube-system        kube-proxy-hphfb                                                     1/1     Running
kube-system        kube-proxy-qzxhr                                                     1/1     Running
kube-system        kube-proxy-rkjds                                                     1/1     Running
kube-system        kube-scheduler-hci-ai-k8-master-01                                  1/1     Running
kube-system        kube-scheduler-hci-ai-k8-master-02                                  1/1     Running
kube-system        kube-scheduler-hci-ai-k8-master-03                                  1/1     Running
kube-system        kubernetes-dashboard-5fcff756f-dmswt                                1/1     Running
kube-system        kubernetes-metrics-scraper-747b4fd5cd-4q4p2                         1/1     Running
kube-system        nginx-proxy-hci-ai-k8-worker-01                                     1/1     Running
kube-system        nginx-proxy-hci-ai-k8-worker-02                                     1/1     Running
kube-system        node-local-dns-2dmjr                                                 1/1     Running
kube-system        node-local-dns-b7xrw                                                 1/1     Running
kube-system        node-local-dns-jrhz2                                                 1/1     Running
kube-system        node-local-dns-jztzs                                                 1/1     Running
kube-system        node-local-dns-wgx84                                                 1/1     Running

```

11. Verify that the Kubernetes setup can access and use the GPUs.

```
./scripts/k8s_verify_gpu.sh
```

Expected sample output:

```

rarvind@deployment-jump:~/deepops$ ./scripts/k8s_verify_gpu.sh
job_name=cluster-gpu-tests
Node found with 3 GPUs
Node found with 3 GPUs
total_gpus=6
Creating/Deleting sandbox Namespace
updating test yaml
downloading containers ...

```

job.batch/cluster-gpu-tests condition met

executing ...

Mon Aug 17 16:02:45 2020

```
+-----+
-----+
| NVIDIA-SMI 440.64.00      Driver Version: 440.64.00      CUDA Version:
10.2      |
|-----+-----+
+-----+
| GPU Name          Persistence-M| Bus-Id        Disp.A | Volatile
Uncorr. ECC |
| Fan  Temp  Perf  Pwr:Usage/Cap|      Memory-Usage | GPU-Util
Compute M. |
|=====+=====+=====+
=====|
|   0   Tesla T4              On   | 00000000:18:00.0 Off  |
0 |
| N/A    38C    P8      10W /  70W |      0MiB / 15109MiB |      0%
Default |
+-----+-----+
+-----+
+-----+
-----+
| Processes:                                     GPU
Memory |
| GPU          PID    Type    Process name                     Usage
|
|=====+=====+=====+
=====|
|   No running processes found
|
+-----+
-----+
Mon Aug 17 16:02:45 2020
+-----+
-----+
| NVIDIA-SMI 440.64.00      Driver Version: 440.64.00      CUDA Version:
10.2      |
|-----+-----+
+-----+
| GPU Name          Persistence-M| Bus-Id        Disp.A | Volatile
Uncorr. ECC |
| Fan  Temp  Perf  Pwr:Usage/Cap|      Memory-Usage | GPU-Util
Compute M. |
|=====+=====+=====+
=====|
```

```

| 0 Tesla T4 On | 00000000:18:00.0 Off |
0 |
| N/A 38C P8 10W / 70W | 0MiB / 15109MiB | 0%
Default |
+-----+-----+
+-----+
+-----+
-----+
| Processes: GPU
Memory |
| GPU PID Type Process name Usage
|
|=====
=====|
| No running processes found
|
+-----+
-----+
Mon Aug 17 16:02:45 2020
+-----+
-----+
| NVIDIA-SMI 440.64.00 Driver Version: 440.64.00 CUDA Version:
10.2 |
|-----+-----+
+-----+
| GPU Name Persistence-M| Bus-Id Disp.A | Volatile
Uncorr. ECC |
| Fan Temp Perf Pwr:Usage/Cap| Memory-Usage | GPU-Util
Compute M. |
|=====+=====+=====
=====|
| 0 Tesla T4 On | 00000000:18:00.0 Off |
0 |
| N/A 38C P8 10W / 70W | 0MiB / 15109MiB | 0%
Default |
+-----+-----+
+-----+
+-----+
-----+
| Processes: GPU
Memory |
| GPU PID Type Process name Usage
|
|=====
=====|
| No running processes found

```



```

|
+-----+
-----+
Mon Aug 17 16:02:45 2020
+-----+
-----+
| NVIDIA-SMI 440.64.00      Driver Version: 440.64.00      CUDA Version:
10.2      |
|-----+-----+
+-----+
| GPU Name          Persistence-M| Bus-Id          Disp.A | Volatile
Uncorr. ECC |
| Fan  Temp  Perf  Pwr:Usage/Cap|      Memory-Usage | GPU-Util
Compute M. |
|=====+=====+=====+
=====|
|   0  Tesla T4              On   | 00000000:18:00.0 Off  |
0 |
| N/A    38C    P8     10W /  70W |      0MiB / 15109MiB |      0%
Default |
+-----+-----+
+-----+
+-----+
-----+
| Processes:                                GPU
Memory |
| GPU          PID    Type    Process name                     Usage
|
|=====+=====+=====+
=====|
|   No running processes found
|
+-----+
-----+
Mon Aug 17 16:02:45 2020
+-----+
-----+
| NVIDIA-SMI 440.64.00      Driver Version: 440.64.00      CUDA Version:
10.2      |
|-----+-----+
+-----+
| GPU Name          Persistence-M| Bus-Id          Disp.A | Volatile
Uncorr. ECC |
| Fan  Temp  Perf  Pwr:Usage/Cap|      Memory-Usage | GPU-Util
Compute M. |
|=====+=====+=====+
=====+

```

```

=====|
|   0   Tesla T4                On   | 00000000:18:00.0 Off |
0 |
| N/A   38C    P8    10W /  70W |      0MiB / 15109MiB |      0%
Default |
+-----+-----+
+-----+
+-----+
-----+
| Processes:                                     GPU
Memory |
| GPU          PID    Type    Process name                      Usage
|
|=====
=====|
|   No running processes found
|
+-----+
-----+
Mon Aug 17 16:02:45 2020
+-----+
-----+
| NVIDIA-SMI 440.64.00    Driver Version: 440.64.00    CUDA Version:
10.2      |
|-----+-----+
+-----+
| GPU Name          Persistence-M| Bus-Id          Disp.A | Volatile
Uncorr. ECC |
| Fan  Temp  Perf  Pwr:Usage/Cap|      Memory-Usage | GPU-Util
Compute M. |
|=====+=====+=====
=====|
|   0   Tesla T4                On   | 00000000:18:00.0 Off |
0 |
| N/A   38C    P8    10W /  70W |      0MiB / 15109MiB |      0%
Default |
+-----+-----+
+-----+
+-----+
-----+
| Processes:                                     GPU
Memory |
| GPU          PID    Type    Process name                      Usage
|
|=====
=====|

```

```
| No running processes found
|
+-----+
-----+
Number of Nodes: 2
Number of GPUs: 6
6 / 6 GPU Jobs COMPLETED
job.batch "cluster-gpu-tests" deleted
namespace "cluster-gpu-verify" deleted
```

12. Install Helm on the Deployment Jump.

```
./scripts/install_helm.sh
```

13. Remove the taints on the master nodes.

```
kubectl taint nodes --all node-role.kubernetes.io/master-
```

This step is required to run the LoadBalancer pods.

14. Deploy LoadBalancer.

15. Edit the `config/helm/metallb.yml` file and provide a range of IP ddresses in the `Application Network` to be used as LoadBalancer.

```
---
# Default address range matches private network for the virtual cluster
# defined in virtual/.
# You should set this address range based on your site's infrastructure.
configInline:
  address-pools:
    - name: default
      protocol: layer2
      addresses:
        - 172.21.231.130-172.21.231.140#Application Network
controller:
  nodeSelector:
    node-role.kubernetes.io/master: ""
```

16. Run a script to deploy LoadBalancer.

```
./scripts/k8s_deploy_loadbalancer.sh
```

17. Deploy an Ingress Controller.

```
./scripts/k8s_deploy_ingress.sh
```

[Next: Deploy and Configure ONTAP Select in the VMware Virtual Infrastructure \(Automated Deployment\)](#)

Copyright Information

Copyright © 2021 NetApp, Inc. All rights reserved. Printed in the U.S. No part of this document covered by copyright may be reproduced in any form or by any means-graphic, electronic, or mechanical, including photocopying, recording, taping, or storage in an electronic retrieval system-without prior written permission of the copyright owner.

Software derived from copyrighted NetApp material is subject to the following license and disclaimer:

THIS SOFTWARE IS PROVIDED BY NETAPP "AS IS" AND WITHOUT ANY EXPRESS OR IMPLIED WARRANTIES, INCLUDING, BUT NOT LIMITED TO, THE IMPLIED WARRANTIES OF MERCHANTABILITY AND FITNESS FOR A PARTICULAR PURPOSE, WHICH ARE HEREBY DISCLAIMED. IN NO EVENT SHALL NETAPP BE LIABLE FOR ANY DIRECT, INDIRECT, INCIDENTAL, SPECIAL, EXEMPLARY, OR CONSEQUENTIAL DAMAGES (INCLUDING, BUT NOT LIMITED TO, PROCUREMENT OF SUBSTITUTE GOODS OR SERVICES; LOSS OF USE, DATA, OR PROFITS; OR BUSINESS INTERRUPTION) HOWEVER CAUSED AND ON ANY THEORY OF LIABILITY, WHETHER IN CONTRACT, STRICT LIABILITY, OR TORT (INCLUDING NEGLIGENCE OR OTHERWISE) ARISING IN ANY WAY OUT OF THE USE OF THIS SOFTWARE, EVEN IF ADVISED OF THE POSSIBILITY OF SUCH DAMAGE.

NetApp reserves the right to change any products described herein at any time, and without notice. NetApp assumes no responsibility or liability arising from the use of products described herein, except as expressly agreed to in writing by NetApp. The use or purchase of this product does not convey a license under any patent rights, trademark rights, or any other intellectual property rights of NetApp.

The product described in this manual may be protected by one or more U.S. patents, foreign patents, or pending applications.

RESTRICTED RIGHTS LEGEND: Use, duplication, or disclosure by the government is subject to restrictions as set forth in subparagraph (c)(1)(ii) of the Rights in Technical Data and Computer Software clause at DFARS 252.277-7103 (October 1988) and FAR 52-227-19 (June 1987).

Trademark Information

NETAPP, the NETAPP logo, and the marks listed at <http://www.netapp.com/TM> are trademarks of NetApp, Inc. Other company and product names may be trademarks of their respective owners.